

Advanced Shortest Paths: Contraction Hierarchies

Michael Levin

Higher School of Economics

Graph Algorithms
Data Structures and Algorithms

Outline

- 1 Contraction Hierarchies
- 2 Preprocessing
- 3 Witness Search
- 4 Query
- 5 Query Correctness
- 6 Node Ordering

Learning Objectives

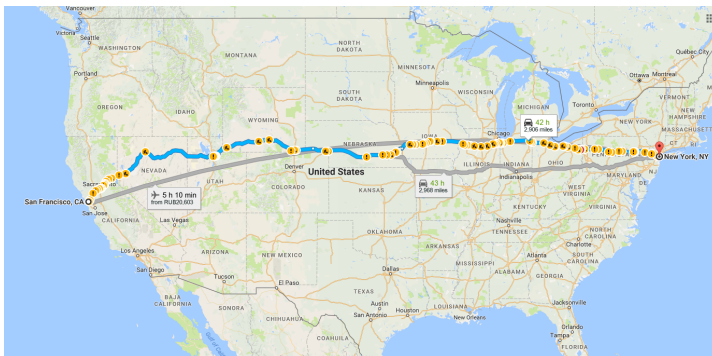
- Bidirectional Dijkstra can be 1000s of times faster than Dijkstra for social networks
- But just 2x speedup for road networks
- This lecture — great speedup for road networks

Highway Hierarchies

- Long-distance trips go through highways

Highway Hierarchies

- Long-distance trips go through highways



Highway Hierarchies

- Long-distance trips go through highways

42 h (2,906 miles)



via I-80 E

40 h without traffic

⚠ This route has tolls

San Francisco, CA

USA

> Get on US-101 S

3 min (0.6 mi)

> Follow I-80 E to Holland Tunnel in Jersey City

42 h (2,903 mi)

> Continue on Holland Tunnel. Drive to Steve Flanders Square in Manhattan, New York

10 min (2.8 mi)

New York, NY

USA

Highway Hierarchies

- Long-distance trips go through highways
- To get from A to B , first merge into a highway, then into a bigger highway, etc., then exit to a highway, then exit to a street, then go to B

Highway Hierarchies

- Long-distance trips go through highways
- To get from A to B , first merge into a highway, then into a bigger highway, etc., then exit to a highway, then exit to a street, then go to B
- Less important roads merge into more important roads

Highway Hierarchies

- Long-distance trips go through highways
- To get from A to B , first merge into a highway, then into a bigger highway, etc., then exit to a highway, then exit to a street, then go to B
- Less important roads merge into more important roads
- Hierarchy of roads

Highway Hierarchies

- There are algorithms based on this idea
- “Highway Hierarchies” and “Transit Node Routing” by Sanders and Schultes
- Millions of times faster than Dijkstra
- Pretty complex
- This lecture — “Contraction Hierarchies”, thousands of times faster than Dijkstra

Node Ordering

Idea behind Contraction Hierarchies - Node Ordering

- Nodes can be ordered by some “importance”
- Importance first increases, then decreases back along any shortest path
- E.g., points where a highway merges into another highway
- Can use bidirectional search

Importance Ideas

Many shortest paths involve important nodes



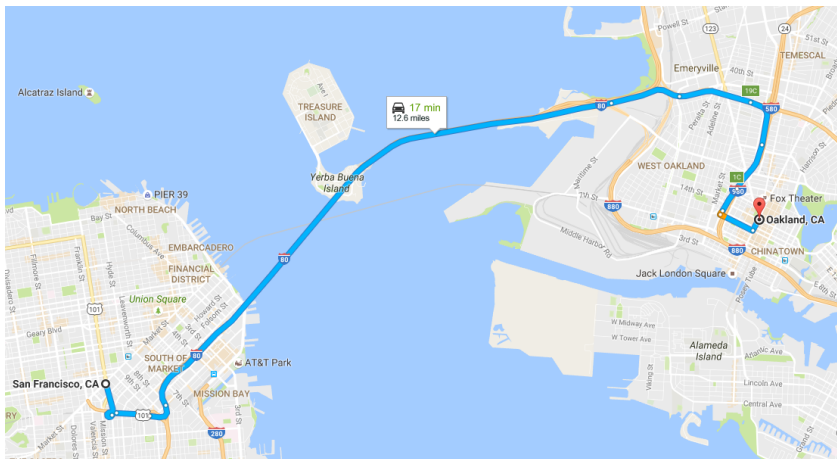
Importance Ideas

Important nodes are spread around



Importance Ideas

Important nodes are sometimes unavoidable



Shortest Paths with Preprocessing

- Preprocess the graph
- Find distance and shortest path in the preprocessed graph
- Reconstruct the shortest path in the initial graph

Question

If we are allowed to do preprocessing, why not just compute all the pairwise distances between the nodes of the initial graph and store them in one table? Then it will only cost $O(1)$ to retrieve the distance between any two nodes.

- ☐ Such preprocessing cannot be done in polynomial time
- ☐ This is impossible at all
- ☒ Such preprocessing can be done, however, for practical applications on the real road networks, such as Europe, USA, or even some big city, it would take too long to compute, and the resulting table would be too big to store

Correct

Indeed, for a graph of USA with 20M nodes and 50M edges, running Dijkstra's algorithm with priority queue from each node would take time proportional to $|V| \cdot |E| \log |V| \sim 20\,000\,000 \cdot 50\,000\,000 \cdot 25 = 25 \cdot 10^{15}$. If we assume a small constant here and roughly 1 billion operations per second, it will still take us about a year to compute the table. Also, storing it would require at least $20\,000\,000 \cdot 20\,000\,000 \cdot 4 = 1.6 \cdot 10^{13}$ bytes or 1.6 Petabytes.


Outline

- 1 Contraction Hierarchies
- 2 Preprocessing
- 3 Witness Search
- 4 Query
- 5 Query Correctness
- 6 Node Ordering


Preprocessing

- Eliminate nodes one by one in some order
- Add **shortcuts** to preserve distances
- Output: augmented graph + node order

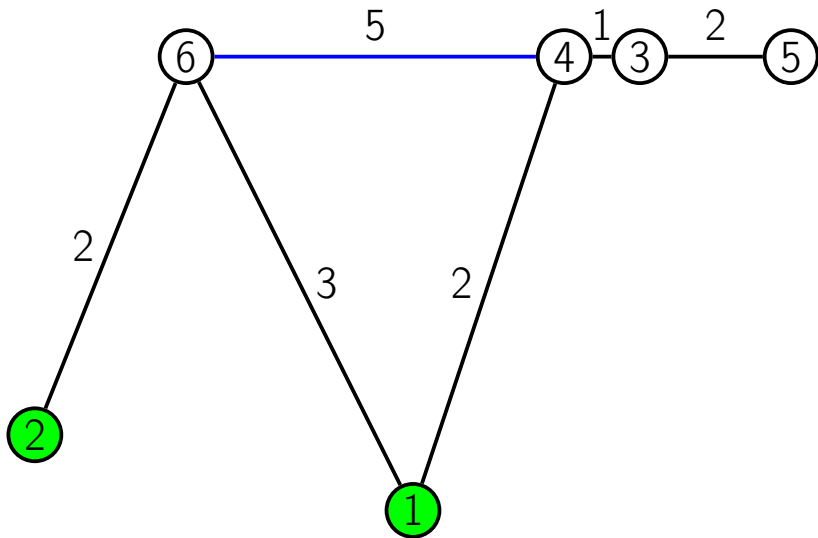
Node Contraction



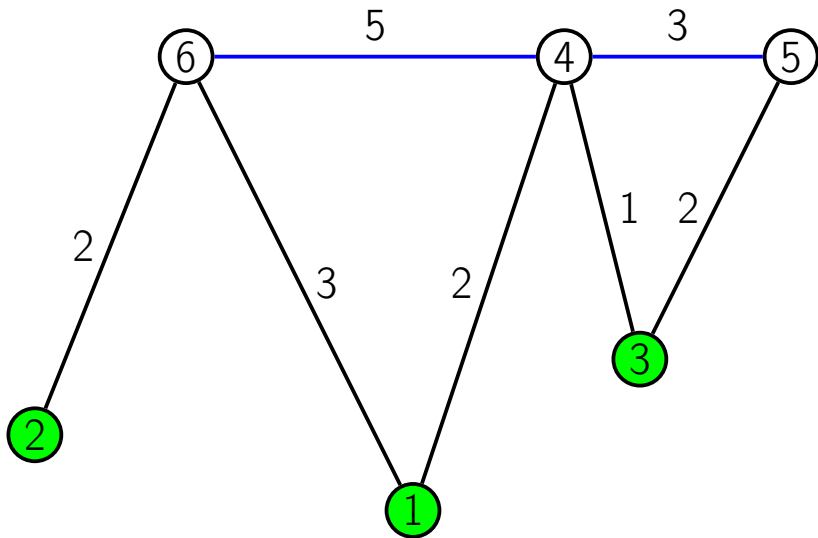
Node Contraction



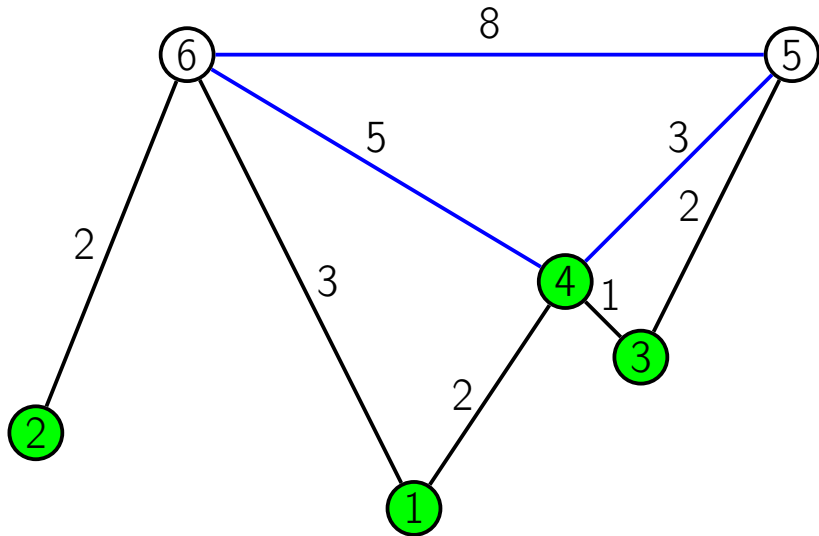
Node Contraction



Node Contraction




Node Contraction



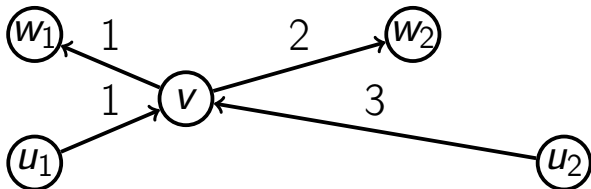
Node Contraction

NO Need to Contract Node 6 as it is the last node



Witness Paths


- Contraction of node v



Witness Paths

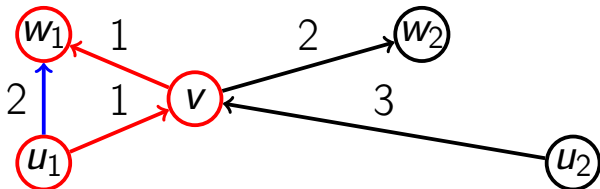
- Contraction of node v
- For every pair of edges (u, v) , (v, w) add a new edge (u, w)

The new path blue is added if we assume that u_1 to v and v to w_1 is the shortest path




Witness Paths

- Contraction of node v
- For every pair of edges (u, v) , (v, w) add a new edge (u, w)
- $\ell(u, w) \leftarrow \ell(u, v) + \ell(v, w)$




Witness Path



- Contraction of node v
- For every pair of edges $(u, \dots, \dots, \dots, w)$, add a new edge (u, w)
- $\ell(u, w) \leftarrow \ell(u, v) + \ell(v, w)$
- But only if there is no **witness path** P_{uw} shorter than $\ell(u, v) + \ell(v, w)$ and

bypassing v If there is a shorter path without including v then we should completely bypass v . These paths are called witness path because these are witnesses that we don't need to add the shortcut and in practice we want to avoid shortcuts as much as possible because lesser the edges the faster the running time



Outline

- 1 Contraction Hierarchies
- 2 Preprocessing
- 3 Witness Search
- 4 Query
- 5 Query Correctness
- 6 Node Ordering

Witness Search

When contracting node v , for any pair of edges (u, v) and (v, w) we want to check whether there is a **witness path** from u to w bypassing v with length at most $\ell(u, v) + \ell(v, w) - 1$ — then there is no need to add a shortcut from u to w .

Definition

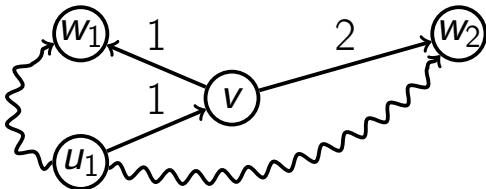
Witness search is the search for a **witness path**.

Definition

If there is an edge (u, v) , call u a predecessor of v . If there is an edge (v, w) , call w a successor of v .

Witness Search

- For each predecessor u_i of v , run Dijkstra from u_i ignoring v
- Essential for good query performance
- Otherwise the augmented graph will be very dense



Witness Search Optimizations

- Stop Dijkstra when distance from the source becomes too big
- Limit the number of hops

Of course we will not find all the witness paths but the speed of pre-processing will increase


Stop Dijkstra

- If $d(u_i, x) > \max_{u, w} (\ell(u, v) + \ell(v, w))$,
there is no witness path going through x

because all the paths going through x will be larger then the $d(u, x)$


Question

How many of the following paths are shortest paths from u to x in the graph below? (The graph is a directed graph with 10 nodes and 15 edges.)



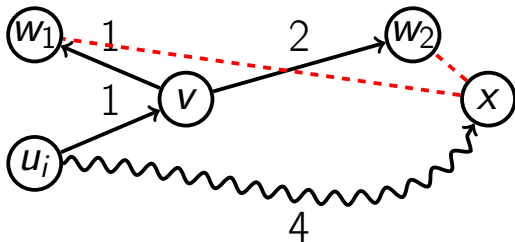
- ☐ shortest path
- ☐ shortest path
- ☐ shortest path
- ☐ shortest path

✓ Answer:
The graph is a directed graph with 10 nodes and 15 edges. The edges and their weights are: $u \rightarrow a$ (3), $u \rightarrow v$ (1), $a \rightarrow v$ (2), $v \rightarrow w_1$ (1), $v \rightarrow w_2$ (2), $v \rightarrow y$ (1), $y \rightarrow x$ (3), $w_1 \rightarrow z$ (1), $w_2 \rightarrow z$ (1), $z \rightarrow c$ (1), $c \rightarrow b$ (1), $b \rightarrow a$ (1). The shortest path from u to x is $u \rightarrow v \rightarrow y \rightarrow x$ with a total weight of 5. The other paths are: $u \rightarrow a \rightarrow v \rightarrow y \rightarrow x$ (6), $u \rightarrow v \rightarrow w_1 \rightarrow z \rightarrow c \rightarrow b \rightarrow a \rightarrow v \rightarrow y \rightarrow x$ (14), $u \rightarrow v \rightarrow w_2 \rightarrow z \rightarrow c \rightarrow b \rightarrow a \rightarrow v \rightarrow y \rightarrow x$ (14), $u \rightarrow v \rightarrow w_1 \rightarrow z \rightarrow c \rightarrow b \rightarrow a \rightarrow v \rightarrow y \rightarrow x$ (14), $u \rightarrow v \rightarrow w_2 \rightarrow z \rightarrow c \rightarrow b \rightarrow a \rightarrow v \rightarrow y \rightarrow x$ (14).



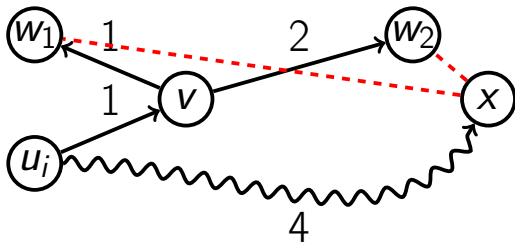
Stop Dijkstra

- If $d(u_i, x) > \max_{u, w} (\ell(u, v) + \ell(v, w))$,
there is no witness path going through x



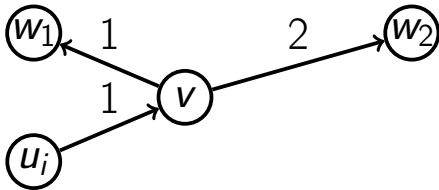
Stop Dijkstra

- If $d(u_i, x) > \max_{u,w}(\ell(u, v) + \ell(v, w))$,
there is no witness path going through x
- Limit the distance by
 $\max_{u,w}(\ell(u, v) + \ell(v, w))$



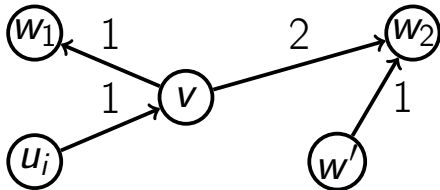
Stop Dijkstra

- Consider any predecessor w' of any successor w of v



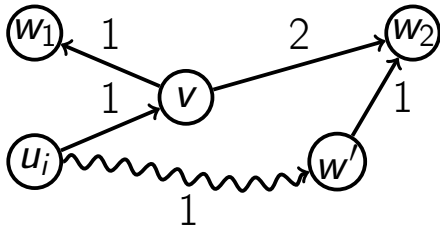
Stop Dijkstra

- Consider any predecessor w' of any successor w of v




Stop Dijkstra

- Consider any predecessor w' of any successor w of v
- If
$$d(u, w') + \ell(w', w) \leq \ell(u, v) + \ell(v, w),$$
there's a witness path



Stop Dijkstra

- Consider any predecessor w' of any successor w of v
- If
$$d(u, w') + \ell(w', w) \leq \ell(u, v) + \ell(v, w),$$
there's a witness path



Stop Dijkstra

■ If

$d(u, w') + \ell(w', w) \leq \ell(u, v) + \ell(v, w),$
there's a witness path


Stop Dijkstra

Very Important Slide. Study Conditions Carefully



If

$$d(u, w') + \ell(w', w) \leq \ell(u, v) + \ell(v, w),$$
there's a witness path




Limit the distance by

$$\max_{u, w} \max_{(w', w)} (\ell(u, v) + \ell(v, w) - \ell(w', w))$$

Question

Which path is the shortest path from node a to node e? (The path is the sequence of nodes visited, and the length is the sum of the edge weights along the path.)



- ☐ a-b-f-g-h-e
☐ a-b-f-g-h-d
☐ a-b-f-g-h-c

Answer

Which path is the shortest path from node a to node e? (The path is the sequence of nodes visited, and the length is the sum of the edge weights along the path.)

Limit the Hops

- Limit the number of “hops” in Dijkstra

Limit the Hops

- Limit the number of “hops” in Dijkstra
- Consider only shortest paths from source with at most k edges

Limit the Hops

- Limit the number of “hops” in Dijkstra
- Consider only shortest paths from source with at most k edges
- If witness path not found, add a shortcut

Limit the Hops

- Limit the number of “hops” in Dijkstra
- Consider only shortest paths from source with at most k edges
- If witness path not found, add a shortcut
- Tradeoff between preprocessing time and augmented graph size

Limit the Hops


- Limit the number of “hops” in Dijkstra
- Consider only shortest paths from source with at most k edges
- If witness path not found, add a shortcut
- Tradeoff between preprocessing time and augmented graph size
- E.g., start with $k = 1$, increase gradually to $k = 5$ in the end

Outline

- 1 Contraction Hierarchies
- 2 Preprocessing
- 3 Witness Search
- 4 Query
- 5 Query Correctness
- 6 Node Ordering

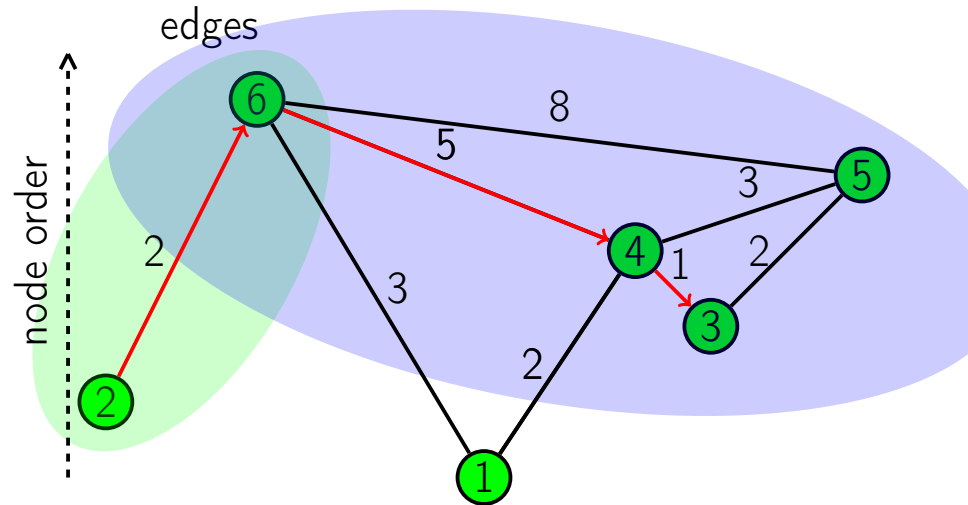
Bidirectional Dijkstra

- Bidirectional Dijkstra using only **upwards**



Bidirectional Dijkstra

- Bidirectional Dijkstra using only **upwards** edges



Bidirectional Dijkstra

- Bidirectional Dijkstra using only upwards edges
- Don't stop when some node was processed both by forward and backward searches
- Stop Dijkstra when the extracted node is already farther than the target

ComputeDistance(s, t, \dots)

```
estimate  $\leftarrow +\infty$ 
Fill dist,  $\text{dist}^R$  with  $+\infty$  for each node
dist[s]  $\leftarrow 0$ ,  $\text{dist}^R[t] \leftarrow 0$ 
proc  $\leftarrow$  empty,  $\text{proc}^R \leftarrow$  empty
while there are nodes to process:
    v  $\leftarrow$  ExtractMin(dist)
    if dist[v]  $\leq$  estimate:
        Process(v, ...)
    if v in  $\text{proc}^R$  and dist[v] +  $\text{dist}^R[v] <$  estimate:
        estimate  $\leftarrow$  dist[v] +  $\text{dist}^R[v]$ 
     $v^R \leftarrow$  ExtractMin( $\text{dist}^R$ )
    Repeat symmetrically for  $v^R$ 
return estimate
```

ComputeDistance(s, t, \dots)

$estimate \leftarrow +\infty$

Fill $dist$, $dist^R$ with $+\infty$ for each node

$dist[s] \leftarrow 0$, $dist^R[t] \leftarrow 0$

$proc \leftarrow \text{empty}$, $proc^R \leftarrow \text{empty}$

while there are nodes to process:

$v \leftarrow \text{ExtractMin}(dist)$

 if $dist[v] \leq estimate$:

 Process(v, \dots)

 if v in $proc^R$ and $dist[v] + dist^R[v] < estimate$:

$estimate \leftarrow dist[v] + dist^R[v]$

$v^R \leftarrow \text{ExtractMin}(dist^R)$

 Repeat symmetrically for v^R

return *estimate*

ComputeDistance(s, t, \dots)

$\text{estimate} \leftarrow +\infty$

Fill dist , dist^R with $+\infty$ for each node

$\text{dist}[s] \leftarrow 0$, $\text{dist}^R[t] \leftarrow 0$

$\text{proc} \leftarrow \text{empty}$, $\text{proc}^R \leftarrow \text{empty}$

while there are nodes to process:

$v \leftarrow \text{ExtractMin}(\text{dist})$

 if $\text{dist}[v] \leq \text{estimate}$:

 Process(v, \dots)

 if v in proc^R and $\text{dist}[v] + \text{dist}^R[v] < \text{estimate}$:

$\text{estimate} \leftarrow \text{dist}[v] + \text{dist}^R[v]$

$v^R \leftarrow \text{ExtractMin}(\text{dist}^R)$

 Repeat symmetrically for v^R

return *estimate*

ComputeDistance(s, t, \dots)

$\text{estimate} \leftarrow +\infty$

Fill dist , dist^R with $+\infty$ for each node

$\text{dist}[s] \leftarrow 0$, $\text{dist}^R[t] \leftarrow 0$

$\text{proc} \leftarrow \text{empty}$, $\text{proc}^R \leftarrow \text{empty}$

while there are nodes to process:

$v \leftarrow \text{ExtractMin}(\text{dist})$

 if $\text{dist}[v] \leq \text{estimate}$:

 Process(v, \dots)

 if v in proc^R and $\text{dist}[v] + \text{dist}^R[v] < \text{estimate}$:

$\text{estimate} \leftarrow \text{dist}[v] + \text{dist}^R[v]$

$v^R \leftarrow \text{ExtractMin}(\text{dist}^R)$

 Repeat symmetrically for v^R

return *estimate*

ComputeDistance(s, t, \dots)

$estimate \leftarrow +\infty$

Fill $dist$, $dist^R$ with $+\infty$ for each node

$dist[s] \leftarrow 0$, $dist^R[t] \leftarrow 0$

$proc \leftarrow \text{empty}$, $proc^R \leftarrow \text{empty}$

while there are nodes to process:

$v \leftarrow \text{ExtractMin}(dist)$

 if $dist[v] \leq estimate$:

 Process(v, \dots)

 if v in $proc^R$ and $dist[v] + dist^R[v] < estimate$:

$estimate \leftarrow dist[v] + dist^R[v]$

$v^R \leftarrow \text{ExtractMin}(dist^R)$

 Repeat symmetrically for v^R

return *estimate*

ComputeDistance(s, t, \dots)

```
estimate  $\leftarrow +\infty$ 
Fill dist,  $\text{dist}^R$  with  $+\infty$  for each node
dist[s]  $\leftarrow 0$ ,  $\text{dist}^R[t] \leftarrow 0$ 
proc  $\leftarrow$  empty,  $\text{proc}^R \leftarrow$  empty
while there are nodes to process:
    v  $\leftarrow$  ExtractMin(dist)
    if dist[v]  $\leq$  estimate:
        Process(v, ...)
    if v in  $\text{proc}^R$  and dist[v] +  $\text{dist}^R[v] <$  estimate:
        estimate  $\leftarrow$  dist[v] +  $\text{dist}^R[v]$ 
     $v^R \leftarrow$  ExtractMin( $\text{dist}^R$ )
    Repeat symmetrically for  $v^R$ 
return estimate
```


ComputeDistance(s, t, \dots)

```
estimate  $\leftarrow +\infty$ 
Fill dist,  $\text{dist}^R$  with  $+\infty$  for each node
dist[s]  $\leftarrow 0$ ,  $\text{dist}^R[t] \leftarrow 0$ 
proc  $\leftarrow$  empty,  $\text{proc}^R \leftarrow$  empty
while there are nodes to process:
    v  $\leftarrow$  ExtractMin(dist)
    if dist[v]  $\leq$  estimate:
        Process(v, ...)
        if v in  $\text{proc}^R$  and dist[v] +  $\text{dist}^R[v] <$  estimate:
            estimate  $\leftarrow$  dist[v] +  $\text{dist}^R[v]$ 
         $v^R \leftarrow$  ExtractMin( $\text{dist}^R$ )
        Repeat symmetrically for  $v^R$ 
return estimate
```

ComputeDistance(s, t, \dots)

estimate $\leftarrow +\infty$

Fill dist, dist^R with $+\infty$ for each node

dist[s] $\leftarrow 0$, $\text{dist}^R[t] \leftarrow 0$

proc \leftarrow empty, $\text{proc}^R \leftarrow$ empty

while there are nodes to process:

$v \leftarrow \text{ExtractMin}(\text{dist})$

 if dist[v] \leq estimate:

 Process(v, \dots)

 if v in proc^R and $\text{dist}[v] + \text{dist}^R[v] < \text{estimate}$:

 estimate $\leftarrow \text{dist}[v] + \text{dist}^R[v]$

$v^R \leftarrow \text{ExtractMin}(\text{dist}^R)$

 Repeat symmetrically for v^R

return estimate

ComputeDistance(s, t, \dots)

$\text{estimate} \leftarrow +\infty$

Fill dist , dist^R with $+\infty$ for each node

$\text{dist}[s] \leftarrow 0$, $\text{dist}^R[t] \leftarrow 0$

$\text{proc} \leftarrow \text{empty}$, $\text{proc}^R \leftarrow \text{empty}$

while there are nodes to process:

$v \leftarrow \text{ExtractMin}(\text{dist})$

 if $\text{dist}[v] \leq \text{estimate}$:

 Process(v, \dots)

 if v in proc^R and $\text{dist}[v] + \text{dist}^R[v] < \text{estimate}$:

$\text{estimate} \leftarrow \text{dist}[v] + \text{dist}^R[v]$

$v^R \leftarrow \text{ExtractMin}(\text{dist}^R)$

 Repeat symmetrically for v^R

return *estimate*

Conclusion

- Preprocessing via nodes contraction

Conclusion

- Preprocessing via nodes contraction
- Query via Bidirectional Dijkstra

Conclusion

- Preprocessing via nodes contraction
- Query via Bidirectional Dijkstra
- Are we done?

Conclusion

- Preprocessing via nodes contraction
- Query via Bidirectional Dijkstra
- Are we done?
- Why is algorithm for query correct?

Because we don't know yet whether this algorithm works correctly or not. We don't know if this estimate is correct. Also it only returns the estimate of the distance from source to target but it doesn't return the actual shortest path between source and target in the initial graph.

- Using the standard bidirectional Dijkstra algorithm we can reconstruct the path in the augmented graph between source and target using the special array in the previous node in the shortest path from source to this node and then same for the backward path

Outline

- 1 Contraction Hierarchies
- 2 Preprocessing
- 3 Witness Search
- 4 Query
- 5 Query Correctness SKIP QUERY CORRECTNESS
- 6 Node Ordering

Augmented Graph

SKIP QUERY CORRECTNESS

Definition

The **augmented graph** $G^+ = (V, E^+)$ is the graph on the same set of vertices V as the initial graph G and an augmented set of edges E^+ that contains all the initial edges E of the graph G along with the shortcuts added at the preprocessing stage.

Distance Preservation

Lemma

The distance $d^+(s, t)$ between any two nodes s and t in the **augmented graph** $G^+ = (V, E^+)$ is equal to the distance $d(s, t)$ between these nodes in the initial graph $G = (V, E)$.

Proof

- Edges are only added to G , so
$$d^+(s, t) \leq d(s, t)$$

Proof

- Edges are only added to G , so $d^+(s, t) \leq d(s, t)$
- For any added shortcut (u, w) , there was a path $u \rightarrow v \rightarrow w$ of length $\ell(u, v) + \ell(v, w) = \ell(u, w)$ before adding this shortcut, so $d^+(s, t)$ can't be less than $d(s, t)$

Proof

- Edges are only added to G , so $d^+(s, t) \leq d(s, t)$
- For any added shortcut (u, w) , there was a path $u \rightarrow v \rightarrow w$ of length $\ell(u, v) + \ell(v, w) = \ell(u, w)$ before adding this shortcut, so $d^+(s, t)$ can't be less than $d(s, t)$
- Thus $d^+(s, t) = d(s, t)$



Definition

The **rank** $r(v)$ of vertex v is the position of v in the node order returned by the preprocessing stage.

Definition

A path $P: v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_k$ in the augmented graph G^+ is called **increasing** if $r(v_1) < r(v_2) < \dots < r(v_k)$. Similarly, P is called **decreasing** if $r(v_1) > r(v_2) > \dots > r(v_k)$.

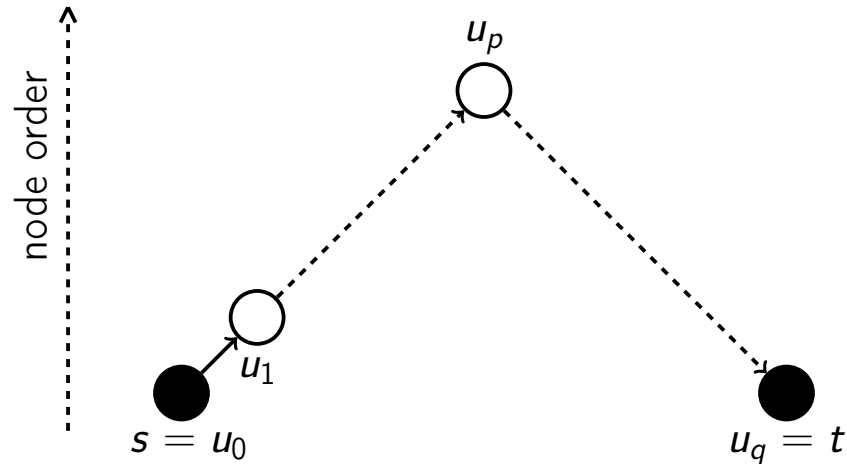
The higher the rank more important it is and later it was contracted

Justification of Bidirectional Search

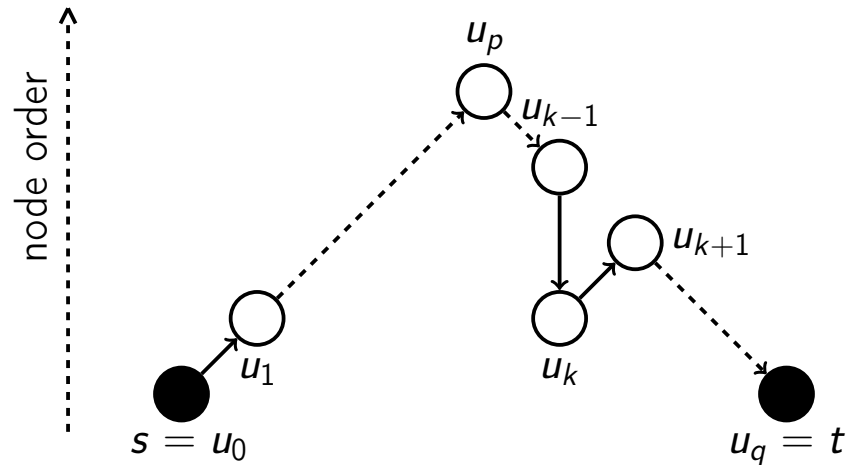
Lemma

For any s and t , the augmented graph $G^+ = (V, E^+)$ contains a shortest path P_{st} such that the subpath P_{sv} is increasing and P_{vt} is decreasing.

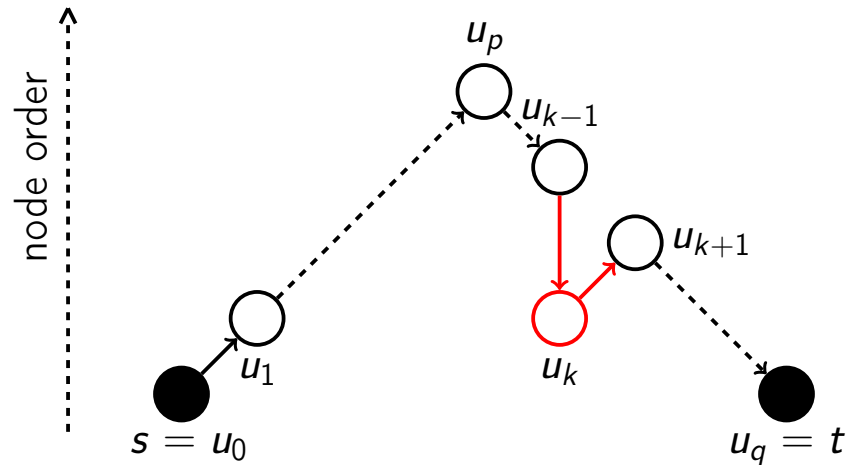
Proof Idea




Proof Idea



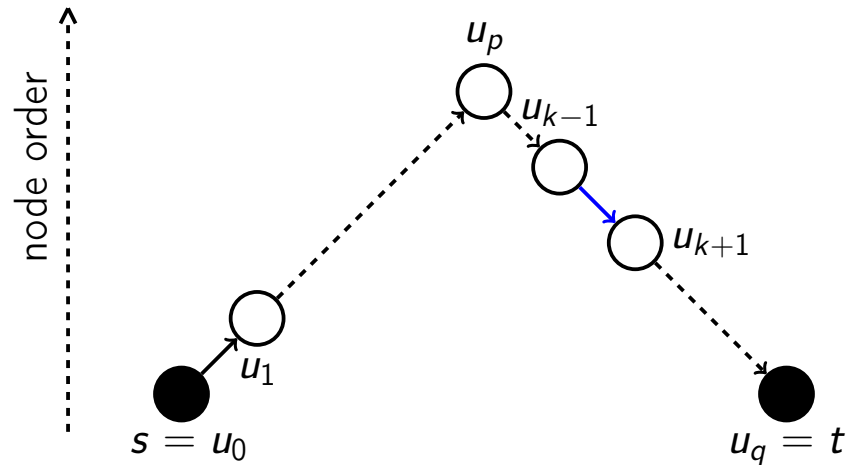
Proof Idea



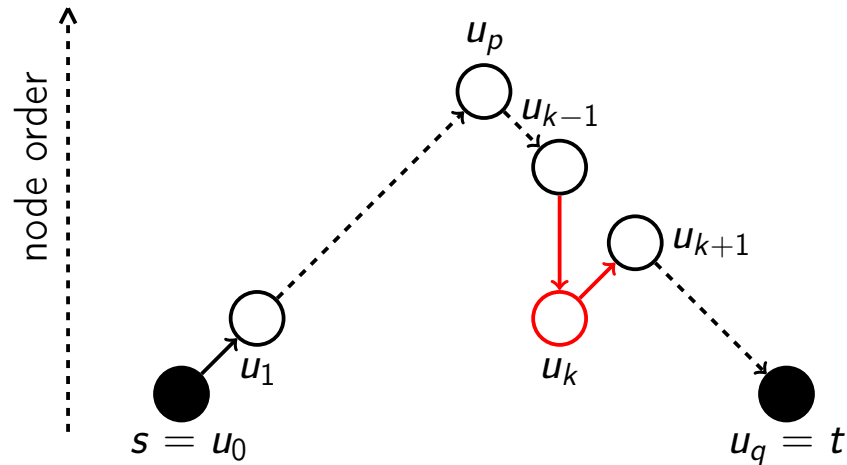
Proof Idea




Proof Idea



Proof Idea



Proof Idea



Proof Idea

Question


Is it possible that after removing u_k and edges $(u_{k-1}, u_k), (u_k, u_{k+1})$ from the path P and replacing them with the witness path $P_{u_{k-1}, u_{k+1}}$ between u_{k-1} and u_{k+1} found while contracting u_k , there will be even more nodes u_i in P are higher than u_i itself?

☒ Yes

☐ No

✓ Correct

It could be, for example, that the witness path between u_{k-1} and u_{k+1} is a path $u_{k-1} \rightarrow v \rightarrow u_{k+1}$ for some v , such that v is higher than both u_{k-1} and u_{k+1} , and then both u_{k-1} and u_{k+1} will become "problematic" in the sense that they are both lower than each of their neighbors in the path P .



Proof

- Assume for the sake of contradiction that no such path P_{st} exists

Proof

- Assume for the sake of contradiction that no such path P_{st} exists
- Then for any shortest path $P: s = u_1 \rightarrow u_2 \rightarrow \dots \rightarrow u_k = t$ there is a node u_i , such that $r(u_{i-1}) > r(u_i) < r(u_{i+1})$ — call it a local minimum

Proof

- For any shortest path P between s and t , denote by $m(P)$ the minimum rank of a local minimum of this path

Proof

- For any shortest path P between s and t , denote by $m(P)$ the minimum rank of a local minimum of this path
- Consider the shortest path P^* with the maximum $m(P)$, consider the local minimum u_k with
$$r(u_{k-1}) > r(u_k) = m(P) < r(u_{k+1})$$

Proof

- If a shortcut (u_{k-1}, u_{k+1}) was added when u_k was contracted, there is a shortest path P' with this shortcut instead of $u_{k-1} \rightarrow u_k \rightarrow u_{k+1}$, and P' doesn't contain u_k , so $m(P') > m(P^*) = r(u_k)$ — contradiction with the choice of P^* with the maximum $m(P)$

Proof

- Otherwise, there was a witness path from u_{k-1} to u_{k+1} comprised by nodes with rank higher than $r(u_k)$ (they were contracted after u_k) — there is a shortest path P'' with this path instead of $u_{k-1} \rightarrow u_k \rightarrow u_{k+1}$, and $m(P'') > m(P^*)$ — contradiction □

Conclusion

- Preprocessing via node contraction

Conclusion

- Preprocessing via node contraction
- Query via Bidirectional Dijkstra in the augmented graph

Conclusion

- Preprocessing via node contraction
- Query via Bidirectional Dijkstra in the augmented graph
- Query correctness is proven

Conclusion

- Preprocessing via node contraction
- Query via Bidirectional Dijkstra in the augmented graph
- Query correctness is proven
- Are we done?

Conclusion

- Preprocessing via node contraction
- Query via Bidirectional Dijkstra in the augmented graph
- Query correctness is proven
- Are we done?
- How to select the node order?

Outline

- 1 Contraction Hierarchies
- 2 Preprocessing
- 3 Witness Search
- 4 Query
- 5 Query Correctness
- 6 Node Ordering

Optimal Node Ordering

- So far our algorithm works for **any** node ordering

Optimal Node Ordering

- So far our algorithm works for **any** node ordering
- However, **preprocessing and query time** depend heavily on it

Optimal Node Ordering

- So far our algorithm works for **any** node ordering
- However, preprocessing and query time depend heavily on it
- Minimize the number of added shortcuts

Optimal Node Ordering

- So far our algorithm works for **any** node ordering
- However, preprocessing and query time depend heavily on it
- Minimize the number of added shortcuts
- Spread the important nodes across the graph

Optimal Node Ordering

- So far our algorithm works for **any** node ordering
- However, **preprocessing and query time** depend heavily on it
- **Minimize the number of added shortcuts**
- **Spread the important nodes** across the graph
- **Minimize the number of edges** in the shortest paths in the augmented graph

Order by Importance

- Introduce a measure of importance

Order by Importance

- Introduce a measure of importance
- Contract the least important node

Order by Importance

- Introduce a measure of importance
- Contract the least important node
- Importance can change after that

Algorithm

- Keep all nodes in a priority queue by decreasing importance

Algorithm

- Keep all nodes in a priority queue by decreasing importance
- On each iteration, extract the least important node

Algorithm

- Keep all nodes in a priority queue by decreasing importance
- On each iteration, extract the least important node
- Recompute its importance

Algorithm

- Keep all nodes in a priority queue by decreasing importance
- On each iteration, extract the least important node
- Recompute its importance
- If it's still minimal (compare with the top of the priority queue), contract the node

Algorithm

- Keep all nodes in a priority queue by decreasing importance
- On each iteration, extract the least important node
- Recompute its importance
- If it's still minimal (compare with the top of the priority queue), contract the node
- Otherwise, put it back into priority queue with new priority

Eventual Stopping

- If we don't contract a node, we update its importance
- After at most $\|V\|$ attempts all nodes have updated importance
- The node with the minimum updated importance will be contracted after that

Importance criteria

- Edge difference Related to the number of shortcuts added and minimizing the number of edges in the augmented graph
- Number of contracted neighbors Related to spreading of nodes across the graph
- Shortcut cover Related to how unavoidable the node is e.g. important node between SF and Oakland
- Node level

Edge Difference

- Want to minimize the number of edges in the augmented graph
- Number of added shortcuts $s(v)$, incoming degree $in(v)$, outgoing degree $out(v)$
- Edge difference
$$ed(v) = s(v) - in(v) - out(v)$$
- Number of edges increases by $ed(v)$ after contracting v
- Contract node with small $ed(v)$

Contracted Neighbors

- Want to spread contracted nodes across the graph
- Contract a node with small number of already contracted neighbors $cn(v)$

Shortcut Cover

- Want to contract important nodes late
- **Shortcut cover** $sc(v)$ — the number of neighbors w of v such that we have to shortcut to or from w after contracting v
- If shortcut cover is big, many nodes “depend” on v
- Contract a node with small $sc(v)$

Node Level

- Node level $L(v)$ is an upper bound on the number of edges in the shortest path from any s to v in the augmented graph
- Initially, $L(v) \leftarrow 0$
- After contracting node v , for neighbors u of v do $L(u) \leftarrow \max(L(u), L(v) + 1)$
- Contract a node with small $L(v)$

Importance

- Use importance

$$I(v) = ed(v) + cn(v) + sc(v) + L(v)$$

- You can play with weights of those 4 quantities in $I(v)$ and see how preprocessing time and query time change
- Each of the 4 quantities is necessary for fast preprocessing/queries
- Find a way to compute them efficiently at any stage of the preprocessing

Comparison with Dijkstra

- On a graph of Europe with 18M nodes, on random pairs of vertices Dijkstra works for 4.365s on average
- On the same graph and same random pairs, with the best set of heuristics Contraction Hierarchies work for 0.18ms on average — almost 25000 times faster!

Conclusion

- Preprocess by contracting nodes ordered *approximately* by importance
- Query by Bidirectional Dijkstra on the augmented graph
- Importance function is heuristic, but works well on road network graphs
- 1000s of times faster than Dijkstra
- Compete on the forums whose solution is the fastest!